

Data processing for IoT


Sandor Markon
Kobe Institute of Computing

Brief self-introduction

- * Born in Budapest, got Dipl. Ing. degree in Electrical Engineering
- * 6 years in electromagnetic field analysis for a Hungarian manufacturer, incl. 2 years in Japan
- * 29 years in computer applications for a Japanese manufacturer, PhD from Kyoto University
- * 12 years teaching at Kobe Institute of Computing
- * Besides: 2.5 companies, visiting researcher at hospital...

What to do with your data in the cloud?

- * Data:
capture - send - collect - **now what???**
- * Convert “data” into “information”
- * Convert “information” into “knowledge”
- * Convert “knowledge” into “action”




Temporal data


Time series: evidence of some process going on
Usual task: **prediction**

What to do with time series?

- * Clean up the dirty data:
 - noise
 - outliers
 - missing data
- * Visualize it
- * Predict the future



Data with outliers




Data with outliers

The spikes are not real data

Why can we predict?

- * Data is generated by some process
- * Usual assumption: deterministic + noise



ARIMA model

$$X_t - \alpha_1 X_{t-1} - \dots - \alpha_{p'} X_{t-p'} = \varepsilon_t + \theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q}$$

- * The measurement at time t depends on previous instances $t-1, t-2, \dots, t-p$
- * It also depends on the “noise” at t and at $t-1, t-2, \dots, t-q$

Time series in Python

- * Numerics: **numpy**, **scipy**

- * Plotting: **matplotlib**

- * Time series: **pandas**


<http://pandas-docs.github.io/pandas-docs-travis/>

- * Prediction (statistical): **pyflux**

<http://www.pyflux.com>


- * Prediction (AI): **FB Prophet**

<https://arnesund.com/2017/02/26/using-facebook-prophet-forecasting-library-to-predict-the-weather/>



ARIMA prediction

We get confidence intervals too...




AI prediction

Generalized model beyond periodicity and trend
(but not very good in this case...)



Healthy



Diseased



Image data

Photos, videos: evidence of some status
Usual task: **classification**

Images + Python + AI




- * Basics: OpenCV
- * Image features: Mahotas
- * Classification: milk

(or scikit-learn: <http://scikit-learn.org/stable/>)



Simple AI for images

- * Generate numerical feature vectors
- * Select part of the images as a training set:
 N_p positive samples
 N_n negative samples
- * Train a classifier with the training set
- * Test it with the rest of the data

negatives






Features




positives


negatives



Features




positives



Neural network


- * Feed features to input layer
- * Calculate hidden and output layer activations through connection weights
- * Compare output with known correct output
- * Adjust weights until output becomes correct for all input samples



Classification


- * After training, the neural network classifies correctly the training samples
- * Hopefully it will also correctly classify unknown data
- * Some common problems (and many others...):
 - * overfitting (can classify only the training set)
 - * poor flexibility (too simple for the task)
- * Illustration: with curve fitting

Training data




Find a good match for the points

Overfitting




Too generic function: cannot find the hidden rule

Underpowered




Too simple function: cannot follow the data

Best fit



2nd order polynomial matches the hidden rule

Best fit, new data



A good fit will match the new data too

Some hints for neural networks

- * Prepare **LOTS** of data!
- * Choose the right network topology (multi-layer, feed-forward; recurrent; etc.)
- * Choose the right complexity
- * Use regularization to fight over-fitting
- * Try self-organizing feature processing